




Jess Lee ++ for The DEV Team

Posted on May 7



215



29



23



22



34

Join the Gemma 4 Challenge: \$3,000 prize pool for TEN winners!

[#devchallenge](#) [#gemmachallenge](#) [#gemma](#)

Announcing the Gemma 4 Challenge

Local AI is having a moment, and we want you to be part of it!

Running through **May 24**, the [Gemma 4 Challenge](#) invites you to explore open models. With the release of Gemma 4, Google's most capable open model family yet, we now have access to native multimodal capabilities, advanced reasoning, a 128K context window, and models that range from running on a Raspberry Pi, to phones, to powering large-scale deployments.

Whether you love to build or love to write, there's a prompt for you and a **\$3,000 prize pool** up for grabs!

Read on to learn more.

Our Prompts

Build With Gemma 4

Your mandate is to **build something useful or creative with any Gemma 4 model**. The scope is wide open — you can build anything from an IoT integration to a multimodal tool to a long-context reasoning app. What matters is that Gemma 4 is doing real work at the heart of your project.

[Build With Gemma 4 Submission Template](#)

The most compelling submissions will make a clear case for *why* you chose the model you did and what that model unlocked.

Write About Gemma 4

Your mandate is to **publish a post about Gemma 4 that educates, inspires, or sparks curiosity**. There's no single right format — what matters is that your post offers something genuine and useful to the community.

Not sure what to write about? Here are some ideas:

- **How-to guide:** Walk through setting up and running a Gemma 4 model locally, fine-tuning it for a specific task, or integrating it into a real project
- **Comparison piece:** Break down the three Gemma 4 model variants and help readers decide which one is right for their use case
- **Personal essay or opinion piece:** Share your experience building with Gemma 4, or make a case for something — what does a model this capable running locally mean for the future of AI?
- **Deep technical breakdown:** Explore a specific capability like multimodal input, the 128K context window, or reasoning mode

[Write About Gemma 4 Submission Template](#)

Need inspiration? Check out how Google's own team [fine-tuned Gemma 4 with Cloud Run Jobs](#). That's the kind of hands-on, shareable knowledge is exactly what we're looking for.

Note: *If you are primarily showing off a project, please submit to the Build with Gemma 4 prompt instead! Each submission is only allowed to be eligible for one of the two categories.*

Which Model Is Right for You? 🤔

Gemma 4 model family spans three distinct architectures tailored for specific hardware requirements:

- **Small Sizes:** 2B and 4B effective parameter models built for ultra-mobile, edge, and browser deployment (e.g., Pixel).
- **Dense:** A powerful 31B parameter dense model that bridges the gap between server-grade performance and local execution.
- **Mixture-of-Experts:** A highly efficient 26B MoE model designed for high-throughput, advanced reasoning.

No matter which you choose, judges will be looking for *intentional* model selection — show us why your model was the right tool for the job.

Prizes 🏆

Five winners from the “Build with Gemma 4” prompt will receive:

- **\$500 USD** cash prize
- **DEV++** Membership
- Exclusive DEV Badge

Five winners from the “Write About Gemma 4” prompt will receive:

- **\$100 USD** cash prize
- **DEV++** Membership
- Exclusive DEV Badge

All participants with a valid submission will receive a **completion badge** on their DEV profile.

Judging Criteria

Build With Gemma 4 submissions will be evaluated for:

- Intentional and effective use of the chosen Gemma 4 model
- Technical implementation and code quality
- Creativity and originality
- Usability and user experience

Write About Gemma 4 submissions will be evaluated for:

- Clarity and depth of explanation
- Originality of perspective or insight
- Practical value to the community
- Quality of writing

Getting Started With Gemma 4

Gemma 4 is open — you have several ways to get started, including completely free options:

- **Gemini API via [Google AI Studio](#):** Access Gemma 4 through the [Gemini API](#).
- **Run locally (free, no credit card required):** Download any Gemma 4 model directly from [Hugging Face](#) or [Kaggle](#). The E2B model runs on high-end phones — or even a Raspberry Pi 5.
- **OpenRouter (free tier available):** Access Gemma 4 31B via [OpenRouter's free tier](#) — no credit card required.

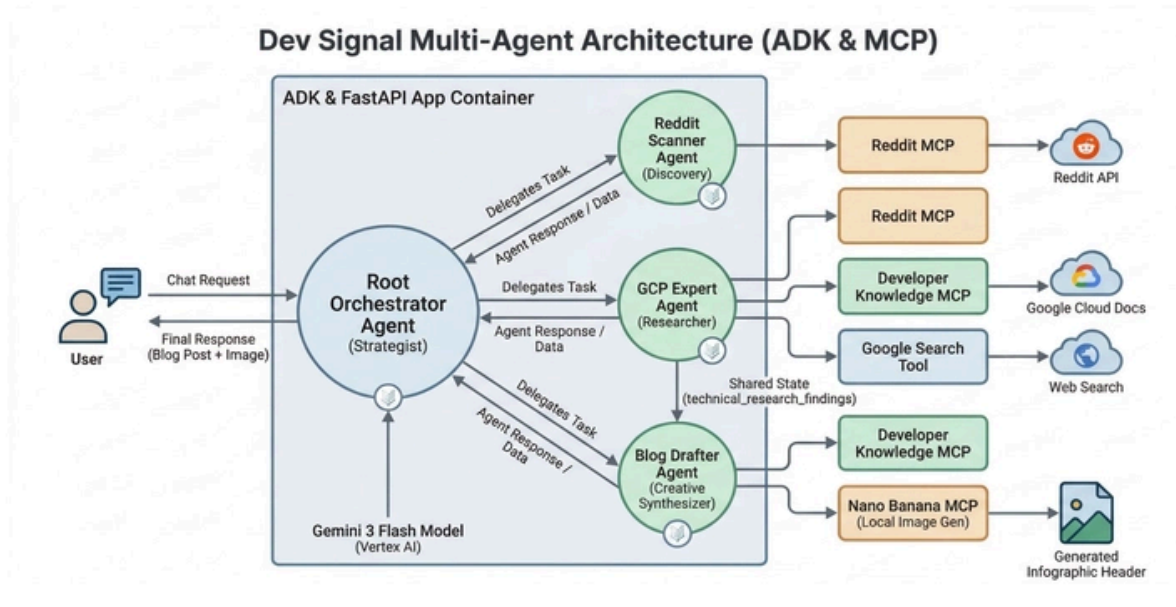
Important Dates

- **May 6:** Gemma 4 Challenge begins!
- **May 24:** Submissions due at 11:59 PM PDT
- **June 4:** Winners Announced

We can't wait to see what you build and write! Questions about the challenge? Drop them in the comments below.

Good luck and happy coding!





Building Capabilities for a Multi-Agent System with Google ADK, MCP, and Cloud Run

My team's mission is to accelerate the developer journey from writing code to running secure AI workloads on Google Cloud. To help developers succeed, we focus on identifying their most pressing questions and building demos that provide straightforward, easy-to-implement solutions.

[Read more →](#)

[Read More](#)

Top comments (59)



FrancisTRDEV (つ〇う〇)つ • May 7



This challenge is either a make or break for me! Will probably do both, but we will see.

@jess If you are doing both prompts, would it be appropriate to do in one post or do you prefer to be separate post. I say this because I was thinking about writing about Gemma 4, but also showcase what I built (not going to spoil it yet).

Regardless, hope to see what the community builds/writes :D



Jess Lee The DEV Team • May 7 • Edited



Hey @francistrdev! I would keep those as two separate submissions. The content should be different though, please don't just duplicate two posts.



FrancisTRDev (つ〇〇〇) • May 7



Sounds good! Thanks! I also sent you an email on other stuff! Noticed you might be busy lol. Imk!



Vikassh. • May 8



Interesting perspective. Curious how others are handling this."



Ben Halpern The DEV Team • May 7



This one is going to be great. I'm most excited about the IOT related use cases, but it all seems cool.



Aashita • May 8



I explored Gemma 4 through the lens of local AI ownership— how moving from cloud-only APIs to capable local models changes the experience for users. One of the most interesting parts for me was intentionally choosing the E4B model instead of defaulting to the largest variant, because accessibility matters just as much as raw capability. Feel free to check out my submission 😊
Excited to read what everyone else is building and writing!



Hemapriya Kanagala • May 7



love this, definitely joining in 😄



Fayaz • May 7



✕ This is a yummy challenge!

Hope to see a lot of really useful projects, and at the same time ***some useless but still cool*** projects!

Love it 🥰

Thanks to Google and the Dev team!



Gábor Mészáros • May 7



✕ Rock & roll



Umar Pathan • May 7



✕ This is a cool challenge. Gemma 4 should lead to some really interesting projects.



Vivek Shetye • May 7



✕ This is exciting. definitely joining this challenge. 👍



HARD IN SOFT OUT • May 7



✕ I want to.. I need to find something I could do with my potato laptop. it's hot even to compile cpp with 20 lines still waiting for +5 mins.



HARD IN SOFT OUT • May 8



✕ ahh, theres solution <3 thank you google & ollama for the cloud <3 amazing!



Paerrin • May 7



✕ Nice! I've already got some good testing with it in my agents memory.

[View full discussion \(59 comments\)](#)


Some comments may only be visible to logged-in visitors. [Sign in](#) to view all comments.

[Code of Conduct](#) • [Report abuse](#)



The DEV Team

PROMOTED



How to prompt Gemini 3.1's new text to speech model

Gemini 3.1 Flash text to speech (TTS) is a new model that you can direct to get the precise audio performance you want. In this blog post I'll share some tips on how to guide the model with prompts, and share some examples of its strengths.

[See more](#)



DEV

The DEV Team

The hardworking team behind DEV ❤️

Want to contribute to open source and help make the DEV community stronger?
The code that powers DEV is called **Forem** and is freely available on GitHub.
You're welcome to jump in!

[Contribute to Forem](#)

More from The DEV Team

Congrats to the OpenClaw Challenge Winners!

[#devchallenge](#) [#openclawchallenge](#) [#openclaw](#) [#ai](#)

Congrats to the Google Cloud NEXT '26 Writing Challenge Winners!

[#devchallenge](#) [#cloudnextchallenge](#) [#googlecloud](#)

Announcing the Winners of the DEV Weekend Challenge: Earth Day Edition 🌍

[#devchallenge](#) [#weekendchallenge](#)

DEV The DEV Team

PROMOTED

...



Explored Web3 but not sure where to start? 🤔

100 Days of Solana is a hands-on challenge series that takes you from curiosity to capability. By the end, you won't just understand Solana, you'll be ready to build something useful with it.

See the latest from [100 Days of Solana on DEV](#).

[Start here](#)